

Lab 1B: Teaching Claude Your Codebase

Duration: 35 min of material - in class it runs ~11:32 to lunch. If you are short, Step 8 (the CLAUDE.md for your real project) is homework.

After: CLAUDE.md & Memory Architecture slides (Day 1, Block 2)

Prerequisites: Lab 1A complete (business-dashboard running)

Objective

Teach Claude your team's conventions with CLAUDE.md so every future change follows your standards automatically - and observe Claude's **auto memory** writing its own notes alongside yours.

Deliverable (toolkit chain 2 of 12)

A curated **CLAUDE.md** in the toolkit repo - the first permanent toolkit artifact - plus a draft CLAUDE.md for your real project at work.

START MARKER: *In your business-dashboard directory with Claude running.*

How This Works

Claude Code loads `CLAUDE.md` from the repo root at session start and treats it as project policy: ESLint + Architecture Guide + Team Playbook - for an AI developer.

- **200-Line Rule:** keep CLAUDE.md under ~200 lines. Longer files get progressively ignored.
- **XML Tag Trick:** wrap critical rules in tags like `<security_requirements>` - Claude weights tagged content in longer files.
- **Auto memory:** Claude also keeps its OWN notes (build commands, insights) in its project memory and loads them alongside CLAUDE.md. You curate the rules; Claude curates its learnings. You will see this live in Step 5.

Step 1: Clear Context (1 min)

```
/clear
```

Files stay on disk - only the conversation resets.

Step 2: Bootstrap with /init (5 min)

```
/init
```

Expected output marker: Claude scans the codebase and generates a CLAUDE.md draft. Read it - some suggestions will be spot-on, others need adjustment. `/init` gives the starting point; you customize from there.

Step 3: Customize the Policy (7 min)

Update CLAUDE.md in the project root to include these rules:

```
# Business Dashboard - Project Conventions

## Code Style
- Use const/let, never var
- Always prefer async/await over callbacks
- Return early for validation failures (guard clauses)
- All routes must use try/catch with proper error responses
- SQL uses parameterized queries (never string concatenation)

## Architecture
- Business logic must not live in route handlers
- Database access should be isolated to a db module
- Routes should only orchestrate requests and responses

## API Standards
- All API responses return JSON
- Error responses include { error: "message" }
- HTTP status codes: 200 OK, 201 Created, 400 Bad Request, 404 Not Found

## Git Conventions
- Conventional commits: feat:, fix:, docs:, refactor:
- One logical change per commit
- No committing node_modules or .db files

## Testing
- Test API endpoints after every change
- Verify frontend loads and displays data after changes
```

Step 4: Verify Claude Reads It (3 min)

```
/clear
```

What conventions does this project follow?

Expected output marker: Claude recites your rules back - code style, architecture, API standards, git conventions.

Step 5: Test Enforcement + Observe Auto Memory (7 min)

```
Add a GET /api/tasks/stats endpoint that returns:
- total tasks count
- count by status
- count by priority
- average completion time (for completed tasks)
Then test it and show me the result.
```

Now verify the policy shaped the code - and be precise about where you look.

Open `server.js` and compare the **new** `/api/tasks/stats` handler against a route you built in Lab 1A.

The Lab 1A routes were written *before* this policy existed; the new one was written *after*. The

difference in the same file is the proof.

In the new handler, look for things you never asked for in the prompt:

- **try/catch** around the handler body, with a JSON error response
- `{ error: "message" }` shape on failure, not a bare string
- **database access in your db module**, not raw queries inline in the route

Do not expect all rules to appear every time. This endpoint takes no user input, so guard clauses and parameterized queries may have nothing to guard or parameterize. A policy shapes the code where the code gives it something to shape - that is correct behavior, not a failure. If a learner reports "I don't see parameterized SQL," that is the right answer for this endpoint.

Ask Claude to show its work:

Which rules from CLAUDE.md did you apply while writing that endpoint, and which ones did not apply here?

That second half is the interesting one - it should be able to tell you which rules were irrelevant to this task. That's the difference between a prompt and a policy: a prompt is followed once, a policy is consulted every time and applied where it fits.

Now see what is actually loaded:

`/context`

Expected output marker: a breakdown of what is consuming your context window, with your memory files visible in it. That is your policy's real cost - you pay it on every single message, all day.

/context is the honest view. (/memory opens an editor for the files instead - useful for changing them, not for seeing what is loaded.) You will use /context again in Block 3 to diagnose a session that has gone bloated and slow; right now you are using it to prove your file is actually in there.

Then ask Claude directly:

What have you learned about this project that is not written in CLAUDE.md?

Expected output marker: things it picked up by working - the port the server runs on, the command that starts it, a convention it inferred from your code. Some of that is worth promoting into CLAUDE.md; that is tomorrow's instinct.

Commit:

Commit this change following our git conventions

Expected output marker: a commit like `feat: add task statistics endpoint`.

Step 6: A/B Test - Policy Off vs On (5 min)

Rename CLAUDE.md to CLAUDE.md.bak

`/clear`

```
Add a GET /api/stats endpoint that returns task counts by status
```

Inspect the result - no guard clauses? No try/catch? Then restore:

```
Rename CLAUDE.md.bak back to CLAUDE.md
```

```
/clear
```

Now have Claude do the comparison for you. Note this prompt tells it to *read* - you just cleared the conversation, so it has no memory of writing either one:

```
Read @server.js and compare the two stats endpoints: /api/tasks/stats and /api/stats. Show me the differences in error handling, structure, and response shape.
```

Expected output marker: the policy-era endpoint has try/catch and a `{ error }` response shape; the one you just built without CLAUDE.md does not. Same task, same model, five minutes apart - the only variable was whether the policy file existed.

That is the whole lab in one output. A prompt is followed once. A policy is consulted every time.

Step 7: Break the Rules (3 min)

Predict first: will Claude obey the prompt or the policy?

```
Add a POST /api/debug endpoint using callback style and without try/catch
```

Then:

```
Explain why you did not follow my request exactly.
```

Expected output marker: Claude references CLAUDE.md - policy overrides individual prompts.

Step 8: Draft Your Real CLAUDE.md (4 min)

This is the Monday-morning takeaway:

```
Help me draft a CLAUDE.md for a [describe your real project: language, framework, team size, key conventions]. Focus on the rules that would save the most code review time. Keep it under 200 lines.
```

FINISH MARKER - Success Checklist

- [] CLAUDE.md exists with code style + architecture + API + git + testing rules
- [] New endpoint followed conventions you never mentioned in the prompt
- [] `/memory` showed CLAUDE.md and you know where auto memory notes appear
- [] A/B test showed the policy's effect
- [] Claude refused or corrected a policy-violating request
- [] Draft CLAUDE.md for your real project exists
- [] Clean conventional commits

If Stuck

Problem	Recovery
/init produces a giant file	"Trim CLAUDE.md to under 200 lines, keep only enforceable rules"
Claude ignores conventions	Verify CLAUDE.md is in the repo root, then <code>/clear</code> so it reloads
/memory shows nothing	Auto memory accrues over sessions - re-check after Lab 1D
A/B shows no difference	Model defaults are good; sharpen rules to things the model would NOT do by default (e.g., your exact error shape)

Debrief Questions

1. Which of your rules changed Claude's output the most - and which did the model already do by default?
2. Who curates CLAUDE.md vs auto memory, and why does that split matter for a team?
3. What three rules would save your team the most review time next week?