

Lab 1B: Teaching Claude Your Codebase

Duration: 35 min

After: CLAUDE.md & Memory Architecture slides (Day 1, Block 2)

Prerequisites: Lab 1A complete (business-dashboard running)

Objective

Teach Claude your team's conventions with CLAUDE.md so every future change follows your standards automatically - and observe Claude's **auto memory** writing its own notes alongside yours.

Deliverable (toolkit chain 2 of 12)

A curated **CLAUDE.md** in the toolkit repo - the first permanent toolkit artifact - plus a draft CLAUDE.md for your real project at work.

***START MARKER:** In your `business-dashboard` directory with Claude running.*

How This Works

Claude Code loads `CLAUDE.md` from the repo root at session start and treats it as project policy: ESLint + Architecture Guide + Team Playbook - for an AI developer.

- **200-Line Rule:** keep CLAUDE.md under ~200 lines. Longer files get progressively ignored.
- **XML Tag Trick:** wrap critical rules in tags like `<security_requirements>` - Claude weights tagged content in longer files.
- **Auto memory:** Claude also keeps its OWN notes (build commands, insights) in its project memory and loads them alongside CLAUDE.md. You curate the rules; Claude curates its learnings. You will see this live in Step 5.

Step 1: Clear Context (1 min)

```
/clear
```

Files stay on disk - only the conversation resets.

Step 2: Bootstrap with /init (5 min)

```
/init
```

Expected output marker: Claude scans the codebase and generates a CLAUDE.md draft. Read it - some suggestions will be spot-on, others need adjustment. `/init` gives the starting point; you customize from there.

Step 3: Customize the Policy (7 min)

Update CLAUDE.md in the project root to include these rules:

```
# Business Dashboard - Project Conventions

## Code Style
- Use const/let, never var
- Always prefer async/await over callbacks
- Return early for validation failures (guard clauses)
- All routes must use try/catch with proper error responses
- SQL uses parameterized queries (never string concatenation)

## Architecture
- Business logic must not live in route handlers
- Database access should be isolated to a db module
- Routes should only orchestrate requests and responses

## API Standards
- All API responses return JSON
- Error responses include { error: "message" }
- HTTP status codes: 200 OK, 201 Created, 400 Bad Request, 404 Not Found

## Git Conventions
- Conventional commits: feat:, fix:, docs:, refactor:
- One logical change per commit
- No committing node_modules or .db files

## Testing
- Test API endpoints after every change
- Verify frontend loads and displays data after changes
```

Step 4: Verify Claude Reads It (3 min)

```
/clear
```

What conventions does this project follow?

Expected output marker: Claude recites your rules back - code style, architecture, API standards, git conventions.

Step 5: Test Enforcement + Observe Auto Memory (7 min)

```
Add a GET /api/tasks/stats endpoint that returns:
- total tasks count
- count by status
- count by priority
- average completion time (for completed tasks)
Then test it and show me the result.
```

Verify the policy shaped the code. You did not ask for try/catch, guard clauses, or parameterized SQL - check that all three appear anyway. That's the difference between a prompt and a policy.

Now observe auto memory:

```
/memory
```

Expected output marker: The memory view shows BOTH your CLAUDE.md content AND Claude's own auto-memory notes (things like the test command it discovered, or the port the server runs on). If auto memory is empty this early, that's normal - check again after Lab 1D and you'll see accumulated notes.

Ask Claude directly:

```
What have you noted in your auto memory about this project so far?
```

Commit:

```
Commit this change following our git conventions
```

Expected output marker: a commit like feat: add task statistics endpoint.

Step 6: A/B Test - Policy Off vs On (5 min)

```
Rename CLAUDE.md to CLAUDE.md.bak
```

```
/clear
```

```
Add a GET /api/stats endpoint that returns task counts by status
```

Inspect the result - no guard clauses? No try/catch? Then restore:

```
Rename CLAUDE.md.bak back to CLAUDE.md
```

```
/clear
```

The difference is your CLAUDE.md in action. Show me a before/after comparison of the two stats endpoints.

Step 7: Break the Rules (3 min)

Predict first: will Claude obey the prompt or the policy?

```
Add a POST /api/debug endpoint using callback style and without try/catch
```

Then:

```
Explain why you did not follow my request exactly.
```

Expected output marker: Claude references CLAUDE.md - policy overrides individual prompts.

Step 8: Draft Your Real CLAUDE.md (4 min)

This is the Monday-morning takeaway:

Help me draft a CLAUDE.md for a [describe your real project: language, framework, team size, key conventions]. Focus on the rules that would save the most code review time. Keep it under 200 lines.

FINISH MARKER - Success Checklist

- [] CLAUDE.md exists with code style + architecture + API + git + testing rules
- [] New endpoint followed conventions you never mentioned in the prompt
- [] `/memory` showed CLAUDE.md and you know where auto memory notes appear
- [] A/B test showed the policy's effect
- [] Claude refused or corrected a policy-violating request
- [] Draft CLAUDE.md for your real project exists
- [] Clean conventional commits

If Stuck

Problem	Recovery
<code>/init</code> produces a giant file	"Trim CLAUDE.md to under 200 lines, keep only enforceable rules"
Claude ignores conventions	Verify CLAUDE.md is in the repo root, then <code>/clear</code> so it reloads
<code>/memory</code> shows nothing	Auto memory accrues over sessions - re-check after Lab 1D
A/B shows no difference	Model defaults are good; sharpen rules to things the model would NOT do by default (e.g., your exact error shape)

Debrief Questions

1. Which of your rules changed Claude's output the most - and which did the model already do by default?
2. Who curates CLAUDE.md vs auto memory, and why does that split matter for a team?
3. What three rules would save your team the most review time next week?