

Lab 3B: Subagents & Agent Teams

Duration: 35 min

After: Subagents & Agent Teams lecture (Day 3, Block 2)

Prerequisites:

- Lab 3A complete (MCP server configured)
- **Phase 4 only - agent teams requires the experimental flag** `CLAUDE_CODE_EXPERIMENTAL_AGENT_TEAMS=1` **set in your environment before you launch** `claude`. Phases 1-3 need nothing extra.

Prerequisite detail: enabling agent teams (experimental flag)

Agent Teams is **disabled by default**. It only starts if the environment variable is set in the shell that launches Claude Code - setting it after launch has no effect.

Windows (PowerShell):

```
```powershell
$env:CLAUDE_CODE_EXPERIMENTAL_AGENT_TEAMS = "1"

claude
```
```

macOS / Linux (bash/zsh):

```
```bash
export CLAUDE_CODE_EXPERIMENTAL_AGENT_TEAMS=1

claude
```
```

Or set it per-project in `.claude/settings.json` so you don't have to remember it:

```
```json
{
 "env": {
 "CLAUDE_CODE_EXPERIMENTAL_AGENT_TEAMS": "1"
 }
}
```
```

If the flag or the feature is unavailable on your machine, that is an expected outcome, not a lab failure. This is an experimental feature and it may be gated, renamed, or unavailable on any given build. Skip to **Phase 4 Step 2 - the parallel-subagent fallback** already written into this lab. The fallback teaches the same coordination lesson using stable subagents, and you are still expected to complete Phase 4 and answer its debrief question.

Objective

Define a custom subagent, run a **nested delegation** (2+ levels deep), add a **grader-style review pass** against a rubric, and launch an **experimental agent team** - with a parallel-subagent fallback if the flag isn't available on your seat.

Deliverable (toolkit chain 10 of 12)

Two subagent definitions in `.claude/agents/` (a builder and a grader) plus a graded, nested multi-agent run you can explain.

START MARKER: `/clear done in business-dashboard.`

Fell behind? Run:

Set up the business-dashboard project with: CLAUDE.md conventions, an Express server with routes in `src/routes/`, and the dashboard-metrics MCP server in `.mcp.json`.

Phase 1: Define Custom Subagents (8 min)

Subagents are markdown files in `.claude/agents/` - isolated instances with their own context, tools, and (optionally) model.

Step 1 - a restricted auditor:

Create a subagent at `.claude/agents/security-auditor.md` in the current project directory:

```
---
name: security-auditor
description: Audit code for security vulnerabilities. Use when asked to "check security",
"audit for vulnerabilities", or "review auth code".
tools: Read, Grep, Glob
---

# Security Auditor
Check for: hardcoded secrets, SQL injection, unvalidated input, eval(), unsafe
deserialization.
Output: CRITICAL / WARNING / INFO findings with file and line references.
```

Step 2 - a grader with a rubric:

Create a subagent at `.claude/agents/grader.md`:

```
---
name: grader
description: Grade completed work against a rubric. Use when asked to "grade", "score
against the rubric", or "run a review pass".
tools: Read, Grep, Glob
---

# Grader
Read the rubric file specified in the task. Score each criterion 0-2 (0=missing, 1=partial,
2=met).
Output a table: criterion | score | evidence (file:line). End with TOTAL: n/10 and a PASS
(>=8) or REWORK verdict. Never fix code - only grade it.
```

Step 3 - verify:

```
/agents
```

Expected output marker: `security-auditor` and `grader` listed alongside built-ins.

Step 4 - first delegation:

Use the `security-auditor` subagent to audit everything in `src/` and save findings to `reports/security.md`

Expected output marker: the subagent runs in its own context, and `reports/security.md` appears with CRITICAL/WARNING/INFO findings. Current Claude Code backgrounds subagents by default; use the subagent panel to inspect status, transcript, tool calls, and usage while the main session remains responsive.

Phase 2: Nested Delegation - 2+ Levels (9 min)

Subagents can spawn subagents. The current default allows up to **three subagent layers below the main conversation** and can be changed with `CLAUDE_CODE_MAX_SUBAGENT_SPAWN_DEPTH` (1 disables nesting). You will run the shallow tree **you -> coordinator -> workers**.

Spawn a subagent acting as a coordinator with this task:
 "Produce `reports/quality-report.md` for the `src/` directory. Do NOT analyze the code yourself - delegate: spawn one subagent to inventory all endpoints and their validation status, and a second subagent to list every function missing error handling. Merge their two outputs into a single report with sections: Endpoints, Error Handling Gaps, Top 3 Risks."

Watch the delegation tree: your session (level 0) -> coordinator (level 1) -> two workers (level 2).

Expected output marker: `reports/quality-report.md` exists with all three sections, and the transcript shows the coordinator spawning its own subagents - not doing the analysis itself.

Show me the delegation tree for that last task: who spawned whom, and what did each level contribute?

***When to nest:** the coordinator holds the MERGE logic; workers hold the DETAIL. Nesting keeps your main context clean - you receive one report, not two raw dumps. Don't nest for simple sequential work; every level multiplies token cost.*

Phase 3: Grader-Style Review Pass (8 min)

The Performance Outcomes pattern: work is not "done" until a separate grader scores it against an explicit rubric.

Step 1 - write the rubric:

Create `rubric.md` with 5 criteria for `src/routes/tasks.js`:

1. All inputs validated with guard clauses
2. try/catch on every route
3. Parameterized SQL only
4. Error responses use `{ error: "message" }`
5. No `console.log` left in production paths

Step 2 - grade:

Use the `grader` subagent to score `src/routes/tasks.js` against `rubric.md`

Expected output marker: a criterion/score/evidence table ending in **TOTAL: n/10** and **PASS** or **REWORK**.

Step 3 - close the loop. If REWORK:

Fix only the criteria the grader scored below 2, then have the grader re-score.

Expected output marker: second grading run scores higher. Builder and grader are SEPARATE contexts - the grader can't be talked into leniency by the builder's reasoning. That separation is the point.

Phase 4: Agent Teams - EXPERIMENTAL (8 min)

EXPERIMENTAL: Agent Teams is disabled by default. Enable it with CLAUDE_CODE_EXPERIMENTAL_AGENT_TEAMS=1 (see the Prerequisites block for PowerShell, bash, and settings.json forms), review the current limitations, and keep the stable subagent fallback ready. (Experimental - confirm it is enabled on your build.)

Step 1 - try to enable. Exit Claude Code first; the variable must be set in the shell **before** launch.

Windows (PowerShell):

```
$env:CLAUDE_CODE_EXPERIMENTAL_AGENT_TEAMS = "1"
claude
```

macOS / Linux:

```
export CLAUDE_CODE_EXPERIMENTAL_AGENT_TEAMS=1
claude
```

Then run:

Create an agent team to add a /api/export/summary endpoint:

- architect: design the endpoint and data shape
- implementer: build it following CLAUDE.md
- reviewer: grade the result against rubric.md

Architect first, then implementer, then reviewer.

Expected output marker: teammates spawn with their own contexts, shared task list, and peer messaging. They normally share the checkout, so the task must partition file ownership; worktrees are not automatic.

Step 2 - FALLBACK (flag or feature unavailable - run this instead, same learning). If the team does not start, do not spend lab time troubleshooting an experimental flag. Confirm you set it before launch, then move straight to this. Completing the fallback fully satisfies Phase 4.

Run this as parallel subagents:

1. One subagent designs the /api/export/summary endpoint and writes the design to reports/design.md
2. When the design exists, a second subagent implements it following CLAUDE.md
3. The grader subagent then scores the implementation against rubric.md

Coordinate the sequence yourself and show me each agent's summary.

Expected output marker: three agents run under your session's coordination - the difference from Teams is WHO coordinates (you/main agent vs the team runtime) and peer-to-peer messaging, which subagents don't have.

Step 3 - decision framework:

| | Single session | Subagents (incl. nested) | Agent Teams |
|--------------|----------------|--------------------------|-------------------------------|
| Coordination | you | main agent | team runtime (peer mailboxes) |

| | Single session | Subagents (incl. nested) | Agent Teams |
|---------|----------------|--------------------------|--|
| Context | one window | isolated per agent | isolated per teammate; shared checkout unless you isolate it |
| Cost | lowest | per-agent multiplier | highest |
| Status | stable | stable | experimental and disabled by default |

FINISH MARKER - Success Checklist

- [] Two subagents defined in `.claude/agents/` with restricted tools
- [] Nested delegation ran: you -> coordinator -> 2 workers (2+ levels)
- [] reports/quality-report.md merged by the coordinator, not by you
- [] Grader scored work against rubric.md with evidence and a verdict
- [] Rework loop raised the score
- [] Agent team launched OR fallback executed, and you can say which and why
- [] You can argue single vs subagents vs teams for a given task

If Stuck

| Problem | Recovery |
|--------------------------------------|---|
| Subagent not found | Files must be in the PROJECT's <code>.claude/agents/</code> ; <code>/clear</code> to reload |
| Coordinator does the work itself | Re-prompt: "You must NOT analyze code directly - delegate to subagents and only merge" |
| Nested spawn refused | Check <code>CLAUDE_CODE_MAX_SUBAGENT_SPAWN_DEPTH</code> ; flatten the design to a shallow coordinator + workers |
| Grader "fixes" the code | Its tools are read-only (Read/Grep/Glob) - if it edits, check the frontmatter tools line |
| Team does not start | Confirm <code>CLAUDE_CODE_EXPERIMENTAL_AGENT_TEAMS=1</code> was set before <code>claude</code> launched (PowerShell <code>\$env :</code> form on Windows), or add it to the <code>env</code> block of <code>.claude/settings.json</code> . If it still doesn't start, the feature is unavailable on your build - run the Step 2 fallback and move on |
| Set the variable but nothing changed | You set it inside an already-running session. Exit, set it in the shell, relaunch <code>claude</code> |
| Token usage exploding | <code>/context</code> ; every agent level multiplies cost - narrow scopes |

Debrief Questions

1. What did nesting buy you that two flat subagents wouldn't have?
2. Why must the grader be a separate agent instead of asking the builder "did you meet the rubric?"

3. Where's the line where you'd move from nested subagents to an agent team - and what would you verify first, given the feature is experimental?